

Creation of a Redshift Cluster

Screenshots of the configuration of the Redshift cluster that you have created:

- Setting up IAM role

IAM > Roles > Create role

Step 1

Select trusted entity

Step 2

Add permissions

Step 3

Name, review, and create

Select trusted entity [Info](#)

Trusted entity type

☒ **AWS service**
Allow AWS services like EC2, Lambda, or others to perform actions in this account.

☐ **AWS account**
Allow entities in other AWS accounts belonging to you or a 3rd party to perform actions in this account.

☐ **Web identity**
Allows users federated by the specified external web identity provider to assume this role to perform actions in this account.

☐ **SAML 2.0 federation**
Allow users federated with SAML 2.0 from a corporate directory to perform actions in this account.

☐ **Custom trust policy**
Create a custom trust policy to enable others to perform actions in this account.

Use case

Allow an AWS service like EC2, Lambda, or others to perform actions in this account.

Common use cases

- ☐ **EC2**
Allows EC2 instances to call AWS services on your behalf.
- ☐ **Lambda**
Allows Lambda functions to call AWS services on your behalf.

Use cases for other AWS services:

Redshift

- ☒ **Redshift - Customizable**
Allows Redshift clusters to call AWS services on your behalf.
- ☐ **Redshift**
Allows Redshift clusters to call AWS services on your behalf.
- ☐ **Redshift - Scheduler**
Allow Redshift Scheduler to call Redshift on your behalf.

Cancel

Next

Add permissions [Info](#)

Permissions policies (Selected 1/808) [Info](#)

Choose one or more policies to attach to your new role.


Filter policies by property or policy name and press enter.

1 match

< 1 > [ⓘ](#)

amazons3fullaccess X

Clear filters

<input checked="" type="checkbox"/>	Policy name ↗	Type	Description
<input checked="" type="checkbox"/>	 AmazonS3FullAccess	AWS m...	Provides full access to all buckets via the AWS Management Console.

▶ Set permissions boundary - optional [Info](#)

Set a permissions boundary to control the maximum permissions this role can have. This is not a common setting, but you can use it to delegate permission management to others.

[IAM](#) > [Roles](#) > [Create role](#)

Step 1
[Select trusted entity](#)

Step 2
[Add permissions](#)

Step 3
Name, review, and create

Name, review, and create

Role details

Role name

Enter a meaningful name to identify this role.

etl-redshift-s3-01

Maximum 64 characters. Use alphanumeric and '+=, @-_' characters.

Description

Add a short explanation for this role.

Allows Redshift clusters to call AWS services on your behalf.

Maximum 1000 characters. Use alphanumeric and '+=, @-_' characters.

Step 1: Select trusted entities


Edit

```
1 {
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Effect": "Allow",
6       "Action": [
7         "sts:AssumeRole"
8       ],
9       "Principal": {
10        "Service": [
11          "redshift.amazonaws.com"
12        ]
13      }
14    }
15  ]
16 }
```

Step 2: Add permissions

Edit

Permissions policy summary

Policy name 	Type	Attached as
AmazonS3FullAccess	AWS managed	Permissions policy

Tags

Add tags - optional [Info](#)

Tags are key-value pairs that you can add to AWS resources to help identify, organize, or search for resources.

No tags associated with the resource.

Add tag

You can add up to 50 more tags.

[Cancel](#)

[Previous](#)

[Create role](#)

- Redshift Cluster Configuration

Create cluster [Info](#)

Cluster configuration

Cluster identifier

This is the unique key that identifies a cluster.

The identifier must be from 1-63 characters. Valid characters are a-z (lowercase only) and - (hyphen).

What are you planning to use this cluster for?

☒ **Production**

Configure for fast and consistent performance at the best price.

☐ **Free trial**

Configure for learning about Amazon Redshift. This configuration is free for a limited time if your organization has never created an Amazon Redshift cluster.

Choose the size of the cluster

☒ **I'll choose**

☐ Help me choose

Node type [Info](#)

Choose a node type that meets your CPU, RAM, storage capacity, and drive type requirements.

Number of nodes

Enter the number of nodes that you need.

Range (1-32)

Configuration summary [Info](#)

dc2.large | 2 nodes

\$360.00/month

Estimated on-demand compute price

Save more than 60% of your costs by purchasing reserved nodes.

[Learn more](#)

320 GB

Total compressed storage

The total storage capacity for the cluster if you deploy the number of nodes that you chose.

Sample data [Info](#)

☐ **Load sample data**

Load sample data to your Redshift cluster to start using the query editor to query data.

Database configurations

Admin user name

Enter a login ID for the admin user of your DB instance.

The name must be 1-128 alphanumeric characters, and it can't be a [reserved word](#).

☐ **Auto generate password**

Amazon Redshift can generate a password for you, or you can specify your own password.

Admin user password

Must be 8-64 characters long. Must contain at least one uppercase letter, one lowercase letter and one number. Can be any printable ASCII character except "/", "", or "@".

☐ Show password

Cluster permissions

i Create an IAM role as the default for this cluster that has the [AmazonRedshiftAllCommandsFullAccess](#) policy attached. This policy includes permissions to run SQL commands to COPY, UNLOAD, and query data with Amazon Redshift. The policy also grants permissions to run SELECT statements for related services, such as Amazon S3, Amazon CloudWatch logs, Amazon SageMaker, and AWS Glue.

Associated IAM roles (1) [Info](#)

Create, associate, or remove an IAM role. You can associate up to 50 IAM roles. You can also choose an IAM role and set it as the default for this cluster.

Set default ▼

Manage IAM roles ▼

< 1 >

<input type="checkbox"/>	IAM roles Info	Status	Role type
<input type="checkbox"/>	etl-redshift-s3-01	Not applied	--

Additional configurations ☒ Use defaults

These configurations are optional, and default settings have been defined to help you get started with your cluster. Turn off "Use defaults" to modify these settings now.

► Network and security [Info](#)

► Database configurations [Info](#)

► Maintenance

► Monitoring

► Backup

Cancel

Create cluster

- Created Redshift cluster
etl-redshift-cluster-1

Actions ▼

Edit

Add partner integration

Query data ▼

General information

Cluster identifier

etl-redshift-cluster-1


Cluster namespace

6339124b-f4dc-4d70-80f7-9423099eb566

Cluster configuration

Production

Status

 Available

Date created

January 23, 2023, 14:29 (UTC+05:30)

Storage used

0.02% (0.07 of 320 GB used)

Multi-AZ

No


Node type

dc2.large


Number of nodes

2


Endpoint

 etl-redshift-cluster-1.cdakre...

JDBC URL

 jdbc:redshift://etl-redshift-cl...

ODBC URL

 Driver={Amazon Redshift (x...

Setting up a database in the Redshift cluster and running queries to create the dimension and fact tables

Queries to create the various dimension and fact tables with appropriate primary and foreign keys:

create schema atm;

✓ creation-i... ✕

✓ dql-tables ✕

+

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

create schema atm;

Run

Save

Schedule

Clear

Query results

Table details

Query

✓ Completed, started on January 24, 2023 at 11:58:30

ELAPSED TIME: 00 m 56 s

```
create table atm.dim_loc
(
    location_id          int primary key,
    location             varchar(30),
    streetname          varchar(30),
    street_number        int,
    zipcode int,
    lat                 decimal(10,3),
    lon                 decimal(10,3)
);
```

Select database [Info](#)

To view schemas, select a database.

dev ▼

Select schema [Info](#)

To view tables, select a schema.

atm ▼

< 1 >

► dim_loc_pkey ...

► dim_loc ...

✓ creation-i... x

✓ dql-tables x

+

```
1 create schema atm;
2
3 create table atm.dim_loc
4 (
5     location_id int primary key,
6     location varchar(30),
7     streetname varchar(30),
8     street_number int,
9     zipcode int,
10    lat decimal(10,3),
11    lon decimal(10,3)
12 );
13
14
15
```

Run

Save

Schedule

Clear

Query results

Table details

Query

✓ Completed, started on January 24, 2023 at 11:59:42
ELAPSED TIME: 00 m 16 s


```
create table atm.dim_atm
(
    atm_id            int primary key,
    atm_number        varchar(20),
    atm_manufacturer  varchar(20),
    atm_location_id   int references atm.dim_loc(location_id)
);
```

Select database [Info](#)

To view schemas, select a database.

dev ▼

Select schema [Info](#)

To view tables, select a schema.

atm ▼

Q Filter tables

< 1 >

▶ dim_atm_pkey ...

▶ dim_loc_pkey ...

▶ dim_atm ...

▶ dim_loc ...

creation-i... X | dql-tables X | +

```
6 location varchar(30),  
7 streetname varchar(30),  
8 street_number int,  
9 zipcode int,  
10 lat decimal(10,3),  
11 lon decimal(10,3)  
12 );  
13  
14 create table atm.dim_atm  
15 (  
16     atm_id int primary key,  
17     atm_number varchar(20),  
18     atm_manufacturer varchar(20),  
19     atm_location_id int references atm.dim_loc(location_id)  
20 );
```

Run

Save

Schedule

Clear

Query results

Table details

Query

Completed, started on January 24, 2023 at 12:01:09
ELAPSED TIME: 00 m 07 s

```
create table atm.dim_date
(
    date_id                int primary key,
    full_date_time         timestamp,
    "year"                 int,
    month                  varchar(20),
    day                    int,
    hour                   int,
    weekday                 varchar(20)
);
```

Select database [Info](#)

To view schemas, select a database.

dev ▼

Select schema [Info](#)

To view tables, select a schema.

atm ▼

Filter tables

<
1
>

- ▶ dim_atm_pkey ...
- ▶ dim_date_pkey ...
- ▶ dim_loc_pkey ...
- ▶ dim_atm ...
- ▶ dim_date ...
- ▶ dim_loc ...

✓ creation-i... ✕
✓ dql-tables ✕
+

```

19   dim_location_id int REFERENCES dim_dim_loc(location_id)
20 );
21
22 create table atm.dim_date
23 (
24     date_id int primary key,
25     full_data_time timestamp,
26     "year" int,
27     month varchar(20),
28     day int,
29     hour int,
30     weekday varchar(20)
31 );
                
```

Run

Save

Schedule

Clear

Query results

Table details

Query

✓ Completed, started on January 24, 2023 at 12:06:24
 ELAPSED TIME: 00 m 15 s

```
create table atm.dim_card_type
(
    card_type_id      int primary key,
    card_type         varchar(30)
);
```

creation-i... x
 dql-tables x
+

Select database [Info](#)

To view schemas, select a database.

dev ▼

Select schema [Info](#)

To view tables, select a schema.

atm ▼

Filter tables

< 1 >

- ▶ dim_atm_pkey ...
- ▶ dim_card_type_pkey ...
- ▶ dim_date_pkey ...
- ▶ dim_loc_pkey ...
- ▶ dim_atm ...
- ▶ dim_card_type ...
- ▶ dim_date ...
- ▶ dim_loc ...

```

25 full_data_time timestamp,
26 "year" int,
27 month varchar(20),
28 day int,
29 hour int,
30 weekday varchar(20)
31 );
32
33 create table atm.dim_card_type
34 (
35     card_type_id int primary key,
36     card_type varchar(30)
37 );
38
39
          
```

Run

Save

Schedule

Clear

Query results
Table details

Query

Completed, started on January 24, 2023 at 12:08:03
ELAPSED TIME: 00 m 06 s

```
create table atm.fact_atm_tran
(
    tran_id                bigint primary key,
    atm_id                 int references atm.dim_atm(atm_id),
    weather_loc_id         int references atm.dim_loc(location_id),
    date_id               int references atm.dim_date(date_id),
    card_type_id          int references atm.dim_card_type(card_type_id),
    atm_status             varchar(20),
    currency              varchar(10),
    service               varchar(20),
    transaction_amount     int,
    message_code          varchar(255),
    message_text          varchar(255),
    rain_3h               decimal(10,3),
    clouds_all            int,
    weather_id            int,
    weather_main           varchar(50),
    weather_description    varchar(255)
);
```

creation-i... x
 dql-tables x
+

Select database [Info](#)

To view schemas, select a database.

dev ▼

Select schema [Info](#)

To view tables, select a schema.

atm ▼

< 1 >

- ▶ dim_atm_pkey ...
- ▶ dim_card_type_pkey ...
- ▶ dim_date_pkey ...
- ▶ dim_loc_pkey ...
- ▶ fact_atm_tran_pkey ...
- ▶ dim_atm ...
- ▶ dim_card_type ...
- ▶ dim_date ...
- ▶ dim_loc ...
- ▶ fact_atm_tran ...

```

36   card_type varchar(30)
37 );
38
39
40 create table atm.fact_atm_tran
41 (
42     tran_id          bigint primary key,
43     atm_id           int references atm.dim_atm(atm_id),
44     weather_loc_id    int references atm.dim_loc(location_id),
45     date_id          int references atm.dim_date(date_id),
46     card_type_id      int references atm.dim_card_type(card_type_id),
47     atm_status        varchar(20),
48     currency          varchar(10),
49     service           varchar(20),
50     transaction_amount int,
51     message_code      varchar(255),
52     message_text       varchar(255),
53     rain_3h           decimal(10,3),
54     clouds_all         int,
55     weather_id        int,
56     weather_main       varchar(50),
57     weather_description varchar(255)
58 );
59 --Card Type

```

Run
Save
Schedule
Clear

Query results
Table details

Query

Completed, started on January 24, 2023 at 12:09:44
ELAPSED TIME: 00 m 08 s

Loading data into a Redshift cluster from Amazon S3 bucket

Queries to copy the data from S3 buckets to the Redshift cluster in the appropriate tables

--Card Type

copy atm.dim_card_type

from 's3://sharads-etl-project/Data/redshift_csv/dim_card_type/part-00000-da182289-b9ce-47d1-8e96-920498350788-c000.csv'

iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'

region 'us-east-1'

CSV;

creation-i... x

dql-tables x

+

```

53  --Location
54  clouds_all      int,
55  weather_id      int,
56  weather_main    varchar(50),
57  weather_description varchar(255)
58 );
59 --Card Type
60 copy atm.dim_card_type
61 from 's3://sharads-etl-project/Data/redshift_csv/dim_card_type/part-00000-da182289-b9ce-47d1-8e96-920498350788-c000.csv'
62 iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'
63 region 'us-east-1'
64 CSV;
65
66
67
68 --Location

```

Run

Save

Schedule

Clear

Query results

Table details

Query 940

Completed, started on January 24, 2023 at 12:17:28

ELAPSED TIME: 00 m 02 s

```
--Location
copy atm.dim_loc
from 's3://sharads-etl-project/Data/redshift_csv/dim_loc/part-00000-ea0bbee5-9124-4aff-95ed-ffc8c7a9d5fc-c000.csv'
iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'
region 'us-east-1'
CSV;
```

creation-i...
dql-tables
+

```

61 --Location
62 copy atm.dim_loc
63 from 's3://sharads-etl-project/Data/redshift_csv/dim_loc/part-00000-ea0bbee5-9124-4aff-95ed-ffc8c7a9d5fc-c000.csv'
64 iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'
65 region 'us-east-1'
66 CSV;
67
68 --Location
69 copy atm.dim_loc
70 from 's3://sharads-etl-project/Data/redshift_csv/dim_loc/part-00000-ea0bbee5-9124-4aff-95ed-ffc8c7a9d5fc-c000.csv'
71 iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'
72 region 'us-east-1'
73 CSV;
74
75
76

```

Run
Save
Schedule
Clear

Query results
Table details

Query [964](#)

Completed, started on January 24, 2023 at 12:18:17
 ELAPSED TIME: 00 m 03 s

```
--ATM
copy atm.dim_atm
from 's3://sharads-etl-project/Data/redshift_csv/dim_atm/part-00000-fae2a521-7e94-4db8-bf64-6f1478f17e06-c000.csv'
iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'
region 'us-east-1'
CSV;
```

creation-i...
dql-tables
+

```

69 copy atm.dim_atm
70 from 's3://sharads-etl-project/Data/redshift_csv/dim_atm/part-00000-fae2a521-7e94-4db8-bf64-6f1478f17e06-c000.csv'
71 iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'
72 region 'us-east-1'
73 CSV;
74
75
76
77 --ATM
78 copy atm.dim_atm
79 from 's3://sharads-etl-project/Data/redshift_csv/dim_atm/part-00000-fae2a521-7e94-4db8-bf64-6f1478f17e06-c000.csv'
80 iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'
81 region 'us-east-1'
82 CSV;
83
84

```

Run
Save
Schedule
Clear

Query results
Table details

Query [1045](#)

Completed, started on January 24, 2023 at 12:19:59
 ELAPSED TIME: 00 m 03 s


```
--Date
copy atm.dim_date
from 's3://sharads-etl-project/Data/redshift_csv/dim_date/part-00000-8e1c5caf-00d4-4f10-9c27-04d48769fb1f-c000.csv'
iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'
region 'us-east-1'
CSV;
```

creation-i...
dql-tables
+

```

79 from 's3://sharads-etl-project/Data/redshift_csv/dim_atm/part-00000-fae2a521-7e94-4db8-bf64-6f1478f17e06-c000.csv'
80 iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'
81 region 'us-east-1'
82 CSV;
83
84
85
86 --Date
87 copy atm.dim_date
88 from 's3://sharads-etl-project/Data/redshift_csv/dim_date/part-00000-8e1c5caf-00d4-4f10-9c27-04d48769fb1f-c000.csv'
89 iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'
90 region 'us-east-1'
91 CSV;
92
93

```

Run
Save
Schedule
Clear

Query results
Table details

Query 1066

Completed
Completed, started on January 24, 2023 at 12:20:53
ELAPSED TIME: 00 m 09 s

--ATM Transaction

copy atm.fact_atm_tran from 's3://sharads-etl-project/Data/redshift_csv/fact_atm_tran1/part-00000-77c6b699-61fb-4ea8-8bc7-d3ae29187fd2-c000.csv'

iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'

region 'us-east-1'

CSV;

creation-i...

dql-tables

+

```

87 copy atm.dim_date
88 from 's3://sharads-etl-project/Data/redshift_csv/dim_date/part-00000-8e1c5caf-00d4-4f10-9c27-04d48769fb1f-c000.csv'
89 iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'
90 region 'us-east-1'
91 CSV;
92
93
94
95 --ATM Transaction
96 copy atm.fact_atm_tran from 's3://sharads-etl-project/Data/redshift_csv/fact_atm_tran1/part-00000-77c6b699-61fb-4ea8-8bc7-d3ae29187fd2-c000.csv'
97 iam_role 'arn:aws:iam::732185110416:role/etl-redshift-s3-01'
98 region 'us-east-1'
99 CSV;

```

Run

Save

Schedule

Clear

Query results

Table details

Query 1103

Completed, started on January 24, 2023 at 12:22:13

ELAPSED TIME: 00 m 20 s